Philosophy and Phenomenological Research
Vol. LXXVI No. 2, March 2008

Assertion, Practical Reason, and Pragmatic Theories of Knowledge

JANET LEVIN
University of Southern California

Defenders of pragmatic theories of knowledge (such as contextualism and sensitive invariantism) argue that these theories, unlike those that invoke a single standard for knowledge, comport with the intuitively compelling thesis that knowledge is the norm of assertion and practical reason. In this paper, I dispute this thesis, and argue that, therefore, the prospects for both "high standard" and "low standard" invariantist theories are better than the pragmatists contend. I also provide a limited defense of the "high standard" approach, and contend that if one abandons the thesis that knowledge is the norm of assertion and practical reason, the most serious arguments against it lose their force.

I. An Argument for Pragmatic Theories of Knowledge

A common feature of contextualist theories of knowledge[1] and the more recent "sensitive invariantist" or "interest-relative" theories (of Jason Stanley and John Hawthorne)[2] is that certain practical considerations have a role in determining whether someone knows that p. On both views, that is, it's possible for A to know that p, and B *not* to know that p, even if A and B have exactly the same evidence for p.

Among the best-known cases designed to prompt intuitions in favor of contextualism are Stewart Cohen's (1999) "airport" scenario and Keith DeRose's (1992) "bank" scenario. In the "airport" scenario,

> Mary and John are at the LA airport contemplating taking a certain flight to New York. They want to know whether the flight has a layover in Chicago. They overhear someone ask a passenger Smith if he knows whether the flight stops in Chicago. Smith looks at the flight itinerary he got from the travel agent and

[1] See, for example, Cohen (1988, 1991), DeRose (1992), Lewis (1996). Some of these theorists take Austin (1946) as a precursor of their views.

[2] See Hawthorne (2004) and Stanley (2005).

responds, ''Yes, I know—it does stop in Chicago.'' It turns out that Mary and John have a very important business contact they have to make at the Chicago airport. Mary says ''How reliable is this itinerary? It could contain a misprint. They could have changed the schedule at the last minute.'' Mary and John agree that Smith doesn't really *know* that the plane will stop in Chicago. They decide to check with the airline agent.

And in the ''bank'' scenario,

> Hannah and her husband are driving home on a Friday afternoon. They plan to stop at the bank on the way home to deposit their paychecks. But as they drive past the bank, they notice that the lines inside are very long, as they often are on Friday afternoons. Thinking that it isn't very important that their paychecks are deposited right away, Hannah says ''I know the bank will be open tomorrow, since I was there just two weeks ago on Saturday morning. So we can deposit them tomorrow morning.'' But then Hannah's husband reminds her that a very important bill comes due Monday, and that they have to have enough money in their account to cover it. He says, ''Banks do change their hours. Are you certain that's not what is going to happen tomorrow?'' Hannah concedes, uttering ''I guess I don't really know that the bank will be open tomorrow.''[3]

For a contextualist, the practical considerations that determine whether S (Smith, John, Mary, Hannah) knows that p have to do with whether or not S has evidence sufficient to rule out those counterpossibilities to p that are salient for the knowledge *ascriber*, the person judging whether or not S has knowledge. Since counterpossibilities to the claim that the plane stops in Chicago have become salient for Mary and John, they can't attribute knowledge that the plane will stop to Smith, while a person for whom these possibilities aren't salient can claim that Smith knows. But since the contexts of attribution are different, they can both be right![4] Similar considerations apply to the case of Hannah and her husband. And, if I've just come from a philosophy class in which we've discussed the ''brain in a vat'' hypothesis, and I'm considering whether you know that you have hands, I have to conclude that you don't—even though the possibility that you're a brain in a vat (and thus don't have hands) has never crossed your mind.

---

[3] This is a slightly condensed version of DeRose's original (1992) bank case, as rendered by Stanley (2005).

[4] This, according to contextualists, is because 'know' shifts its meaning from context to context, and they're thus asserting different propositions.

For a sensitive invariantist, on the other hand, whether S knows that p in a particular situation depends on considerations of what is salient for *S herself*—whether she's thinking (or should be thinking) about certain counterpossibilities—or how important it is for her, given her current situation, to get things right.[5] In the airport case, since it makes little difference to Smith whether the plane stops in Chicago, he counts as knowing that it will, for any ascriber, whereas Mary and John, in contrast, have no such knowledge. And once it becomes clear how important it is for Hannah to deposit her paycheck by Monday, no ascriber should grant her knowledge that the bank will be open on Saturday.[6]

Another important class of cases that both contextualists and sensitive invariantists take to support their views involve so-called "lottery propositions." When thinking about one's chances of winning a lottery (with many tickets sold), one realizes that it's overwhelmingly probable that one will lose. Still, it seems intuitively wrong to claim (or believe) that one knows one has lost the lottery before the winner has been announced—at least when one is concentrating on the arbitrary methods by which winners are determined: *someone* has to win, and there's no reason it won't be me! Yet, even when we own a lottery ticket, we make claims, and plans, that seem to be based on our knowing that we'll lose. Pragmatists of both persuasions suggest that, when the mechanics of the lottery drawing are salient (or important), one *doesn't* know that one's ticket is a loser, while in other situations one *does*. Whether these features of the situation have to be salient (or important) to the ticket holder or the *ascriber* is the point of contention between the two views.[7]

Moreover, as a number of theorists have recently argued, the proposition that one has a losing lottery ticket is not an isolated special case that can be handled by a footnote to one's theories, but one of a substantial class of "lottery propositions" that behave in the same way. Consider, for example, the proposition that one's car hasn't been stolen—when it's parked out of sight on the streets of a city that has a few (arbitrarily chosen) cars stolen every year.[8] One can't really say, it seems, at least when

[5] Sensitive invariantists, in contrast to contextualists, don't propose that the word 'know' changes its meaning from context to context, but rather that features of the subject's situation (what's salient, what's at stake) help to determine whether the (semantically invariant) term can be applied.

[6] The relative importance of "stakes" vs "salience" can differ on different sensitive invariantist views—and can make a difference to the verdict on whether Hannah can be said to know that the bank is open on Saturday *before* her husband points out how important it is that she deposit her check before Monday, and notes that banks sometimes change their hours.

[7] This is, I acknowledge, a simple-minded and un-nuanced sketch of both views, and their differences. It's merely meant to be a sketch.

[8] This example was first given by Jonathan Vogel (1990a) who is neither a contextualist nor a sensitive invariantist, but Hawthorne (2004) discusses these cases at length.

thinking about the arbitrariness of car theft (or when immediate access to one's car assumes tremendous importance), that one *knows* that one's car hasn't been stolen, but most of the time one can be attributed such knowledge free and clear. Another striking example is the proposition that I (if young and healthy, that is) won't have a fatal heart attack in the next week. Even though I assume this in ordinary situations, I'm loath to say (or believe) that I know it when I'm confronted with the statistics about heart attack deaths in my demographic, even if the risks are extremely small.[9] These intuitions are taken to be further evidence for pragmatic theories of knowledge.

However, although both contextualism and sensitive (interest-relative) invariantism match some of our intuitions about whether or not S has knowledge, there are consequences of both views that are more problematic.[10] This seems true as well when the "pragmatic" considerations determine not whether S knows that p, but whether S has justified belief that p.[11] Thus, other philosophers have attempted to take the pragmatic considerations that distinguish these cases to determine not whether S has knowledge, or even justified belief, but whether *S believes that p at all*![12] But an account of knowledge that locates the pragmatic element here has problems of its own.[13]

---

9 See Hawthorne (2004). But see Adler (2005) for a compelling argument that the "car theft" and "heart attack" propositions are significantly different from the (literal) lottery proposition, and that intuitions about whether we can know we've lost the lottery do not carry over to these other cases. This line of argument contributes another way to challenge pragmatic accounts of knowledge, while allowing us nonetheless to have substantial knowledge of the external world.

10 There's a tremendous literature on contextualism (both criticisms and defenses); less on sensitive or interest-relative invariantism, since these are more recent theories. For good examples of each, see Hawthorne (2004) and Stanley on contextualism, and McGrath (2004) and Brueckner (2004) on Hawthorne's sensitive invariantism.

11 See Fantl and McGrath (2002), and, in criticism, Brueckner (forthcoming), and Weatherson (2005).

12 See Weatherson (2005). Indeed there are yet other alternatives. Maybe the pragmatic considerations don't determine whether S believes that p, but whether S can be thought to have *evidence* for p. This is different from claiming that pragmatic considerations determine whether one is justified in believing that p. Under certain conditions, that is, one can take q, r, s, to support p, but under other conditions, the evidential relevance of q, r, s to p may be questioned. For example, reliable testimony supporting p, from a person S believes to be reliable, may standardly count as evidence for p, but in certain high stakes situations, or situations in which one's general faculties are called into doubt, such testimony won't count merely as insufficient evidence for p, but rather won't count as raising the probability of p at all! (See Vogel (1999), note 11, for the suggestion that this is the situation for the skeptic.) I imagine, however, that there are analogous difficulties for *this* "pragmatic" account as well.

13 Detailed by Weatherson himself in his (2005).

One might thus think that, on balance, it's best to embrace a non-pragmatic approach to epistemic evaluation. Why not be, as Hawthorne might put it, an *insensitive* invariantist (or "evidentialist"),[14] and claim that there is one standard for knowledge (and/or justified belief) that should be applied to people, regardless of what's salient or important, either to them or those attempting to ascribe knowledge to them. Insensitive invariantists, of course, must argue for a particular set of (invariant) epistemic standards—either "high," which would make it hard to have knowledge (or justification), or "low" which would make these achievements easier. But pragmatists have argued that neither approach can withstand serious scrutiny, since neither comports with an important and compelling thesis, namely, that *knowledge is the norm of assertion and practical reason*.[15] My primary aim in what follow is to challenge this thesis, and argue that the prospects for either invariantist view are thereby better than the pragmatists contend.[16]

First, I'll consider the prospects for a (relatively) "high standard" approach, on which we don't have much knowledge about the world, and then for a "low standard" approach, which can preserve the bulk of our ordinary knowledge claims. Though I'll offer a tentative defense of a high standard approach, I won't argue that it's clearly the preferable view. I will argue, however, that the intuitive grounds for preferring one to the other look very different once one abandons the thesis that knowledge is the norm of assertion and practical reason.

II. High Standard Theories and Our Ordinary Knowledge Claims

High standard theories (generally)[17] contend that in *no* case do the people in the airport know that the plane will stop in Chicago, given that they've merely checked the (up-to-date) itinerary, and that Hannah doesn't know that the bank will be open on Saturday, given that she's merely checked the hours posted on the door—regardless of the proposition's importance, or the salience of the counterpossibilities to the claim. They also contend that we don't know that our lottery ticket will lose, or that our car, parked out of sight, hasn't been stolen, or that we won't have a fatal heart attack within the week. The reason in

[14] See Conee and Feldman (2004). But, of course, this sort of view was the norm before the advent of contextualist theories.

[15] See, for example, DeRose (2002, p. 182), who claims that "[t]he knowledge account of assertion demands a contextualist theory of knowledge" but continues "and is simply incredible without it."

[16] Indeed, many *invariantists* endorse this claim as well, and use it in support of their particular version of invariantism. I'll discuss this strategy in the final section.

[17] The verdicts in these cases, of course, will depend on the details of the particular high standard view.

all these cases is that there are "real", and not merely "hyperbolic," counterpossibilities to the claims in question which the putative knowers are unable to rule out: things like this *have happened*, it is often argued, and even though the chance of their happening in these cases is small, it cannot be ruled out. Denying knowledge in these cases need not lead to Cartesian skepticism if a principled distinction can be made between "real" (if remote) versus "hyperbolic" (or "Cartesian") counterpossibilities—and there are a number of promising suggestions to this effect in the literature.[18] But even if a distinction can be made that stops the slide into radical skepticism, a serious worry remains.

The worry is that, given the closure of knowledge under known a priori entailment, this lack of knowledge ramifies to an alarming degree. If we don't know that the lottery ticket we just bought is a loser, and recognize that this proposition is entailed by the proposition that we won't have enough money for a trip to Africa next year, then we can't know that we won't have enough money for the trip. Similarly, if we don't know that our car hasn't been stolen, and we recognize that the claim that our car is parked on Elm Street entails that it hasn't been stolen, then we can't know where our car is parked either. But it seems that we *do* know these ordinary propositions; thus, given closure, we have to be able to affirm that we know the propositions they entail—and this is just what high standard invariantist views deny.[19]

---

[18] For one thing, we've observed that people just like ourselves (given all available evidence) have won lotteries, had cars stolen, had fatal heart attacks, and gotten misinformation about schedules, but we've never observed (nor have evidence for the existence of) brains in vats, or victims of any other sort of wholesale deception. For another, the methods by which lottery winners, car theft victims, and (given all available evidence) heart attack victims in my demographic are selected is arbitrary; there's nothing (given all available evidence) that differentiates me from the rest of the pack. Indeed. this (at least arguably) makes these possibilities different from the ("real") possibility that I'm hallucinating, or seeing a mirage, since in cases where those possibilities are real, rather than hyberbolic, there's usually some (discernable) feature of one's situation that calls for caution. See Vogel (1990b, 2004a, 2005) for some interesting ways of distinguishing "hyperbolic" from "real" possibilities.

[19] One response to this problem is to deny that knowledge is closed under known a priori entailment. If we deny closure, then—at least on *some* non-pragmatic views—we can affirm both that we know that we won't be going to Africa, or that our car is parked on Elm Street, while denying that we know that we'll lose the lottery, or that our car hasn't been stolen. Many theorists, however, argue fiercely for the intuitive plausibility of closure, and find any view that denies it to have serious intuitive costs.

I myself am not so worried about denying closure. As Hawthorne (2004, p. 39) puts it, denying closure "interacts disastrously" with the idea that knowledge is the norm of assertion, for if closure fails, we'd sometimes be in the position of asserting 'p', and also 'p⊃q', but refuse to assert q—and also refuse to take back our assertion that p⊃q. But, if knowledge is not the norm of assertion or practical reason, there are fewer practical consequences if closure fails, and thus—perhaps—it will not be as counterintuitive as Hawthorne contends. However, I won't pursue this question here.

But why is it so unacceptable to think that we *don't* know that our car is parked on Elm Street, or that we won't be going to Africa next year? There will still be quite a few ordinary propositions that we, at least arguably, *can* be said to know: some based on perception, memory, or reliable testimony—at least in certain circumstances, and when there is no defeating background information—and some arrived at by rational reflection.[20] And there will be some propositions about the future and unobserved events that we, at least arguably, may be able to know as well: if we see the flower pot get knocked off the table, we can know that soon it will hit the ground; if we see the baby bite into an unripe persimmon, we can know that soon it will pucker its lips.[21] Why do we need to have knowledge about whether the bank is open on Saturday, whether the plane stops in Chicago, whether our car is parked on Elm Street, and various other ordinary propositions that we standardly and unreflectively claim that we know?

---

[20] One may wonder whether a high standard theorist can affirm even this much "ordinary" knowledge; after all, it sometimes happens that peoples' perceptions are misleading, or memories unreliable; how can I rule out the possibility that this isn't now happening with me? What a high standard theorist must argue is that there is a principled way (and, if one is an internalist, a way that's *accessible* to the subject) of differentiating cases in which counterpossibilities to one's claims are real versus "hyperbolic." For example, one might restrict the cases in which perception is considered trustworthy to situations in which certain (discernable) environmental and perceptual standards are met (e.g., the light is on, I'm wearing my glasses, I haven't taken drugs), and the cases in which memory is trustworthy to situations in which I'm sufficiently young, and am suffering from no memory-impairing disease—and, perhaps, that the remembered situation isn't fraught with emotion. In no case will this be an easy task—but there is some reason to think that it can be accomplished. Even so, one might worry, even in these very circumscribed situations, there's *some* (if very slight) "real" chance that I'm mistaken about what I observe, or remember, or even introspect. And if I can't know I've lost the lottery, even though my chances of winning are extremely low, there are very few ordinary propositions about the world that I can know. (Thanks to an anonymous referee for pressing me on this point.)

It may be that, to preserve any but the most basic knowledge about the world, one has to treat propositions like the lottery proposition, for which we have evidence based *only* on probabilistic considerations, as special cases, and suggest that other considerations permit us to take perception, memory, reflection, etc.—at least under certain circumstances—to be sources of knowledge, even though there is *some* probability that claims made on the basis of these methods are false. Even so, the standards for knowledge will be rather demanding, and prevent many ordinary claims about the world from counting as knowledge. I'll provide a (modest) defense of this sort of high standard view in the final section—but, once again, my primary aim is to challenge the thesis that knowledge is the norm of assertion and practical reason.

[21] Again, one can argue that there is reason to think that high standard theories don't have to embrace complete inductive skepticism—but that, even if they do, it won't be so intuitively devastating if the main thesis of this paper is correct.

This is where the question of norms involving knowledge looms large. We can't give up on having knowledge of these "ordinary" propositions so quickly, many theorists contend, because knowledge is the state to which belief "aspires";[22] knowledge is the norm of assertion; knowledge is the norm of practical deliberation. Appeal to these principles has taken on a quasi-moral fervor:[23] if we relinquish what we've ordinarily taken to be knowledge, then *the normative appropriateness of many of our assertions and actions will be undermined*.[24]

But what if belief *doesn't* aspire to knowledge, and knowledge is *not* the norm of assertion, and what if having knowledge (in contrast to various other epistemic and non-epistemic notions) has little to do with how we are rationally required to deliberate and act? What if, instead, it is the *norms for assertion and practical reasoning themselves* that shift, according to the subjects' interests, values, stakes, or other pragmatic considerations? Then presumably, there wouldn't be onerous *practical* consequences of our not knowing various "ordinary" propositions about the world, since, at least in some circumstances, we could perfectly well assert, or use as premises in our practical reasoning, propositions that we don't know.[25]

We could also offer an alternate explanation of what many theorists take to be a phenomenon that supports "pragmatic" accounts of knowledge, namely, that though reflection on "lottery" or "car theft" possibilities seems to undermine our knowledge of the ordinary propositions that entail their negations, we forget about these possibilities when ordinary concerns once again crowd in. We need not say, with the pragmatists, that (for one reason or another) we *don't* know these ordinary propositions when we (or those ascribing knowledge to us) are reflecting on stakes or counterpossibilities, and *do* know them when we're not, but can embrace the (Humean) view that we recognize, upon reflection, that we don't, and never did, know these propositions—but when ensconced at the backgammon table, or engaged in the other

---

22 Williamson (2000), p. 1.

23 See Hawthorne (2004), Stanley (2004 and forthcoming).

24 As Hawthorne puts it (2004, p. 132), "A real trouble spot for skepticism [and presumably, any high standard view] arises from the normative connections between knowledge, assertion, and practical reasoning."

25 If knowledge is not taken to be the norm of assertion and practical reasoning, it will also be easier to defend more liberal, but still invariant and closure-respecting, accounts of knowledge. I discuss this prospect in the final section.

routines of everyday life, it just doesn't matter. I'll argue, in what follows, that this is a more reasonable view of our epistemic (and practical) situation than it may at first seem.

III. Knowledge and the Norms of Assertion

What exactly does it mean to say that knowledge is the norm of assertion? Stanley glosses this as the claim, taken from Unger (1974), that when one asserts that p, one "represents oneself as knowing" p. If this is intended to characterize a norm of assertion, then presumably one can be rationally criticized for asserting p if one does not in fact know that p, since this would be a matter of false representation. My question is whether we do in fact think that proper assertion should be restricted to the propositions one knows, and also whether we do indeed take people who assert that p to be representing themselves as knowing that p.

Some examples, at least on the face of it, suggest that the answer to both questions is "no." Certainly, asserting that p often has the goal of representing oneself as knowing that p, but people make assertions for all sorts of other reasons as well: to establish authority, to demonstrate conviction, to encourage others to believe, to commit themselves publicly to a belief about which they themselves had been (perhaps irrationally) wavering. And assertions are often understood by their audience as doing just these things. Moreover, even when assertions are made strictly for the purpose of conveying information, or giving one's audience the right to assert those things themselves, in many cases they'll be normatively appropriate (and sometimes they'll be normatively required) as long as they express well-justified, sincere beliefs.

For example, a US government official, upon learning that a cow across the Canadian border has tested positively for BSE, may go to the nearest McDonalds, buy a hamburger, and assert (truly), "This burger is perfectly safe," supporting this claim by citing differences in the feed composition, testing, and meat grinding procedures for animals in the US and Canada. Even if we take the official to be sincere, honest and thorough in his research,[26] I suspect we'd think that his chances of being correct are—oh, almost as good as my chances of losing a hundred million ticket lottery.[27] And certainly,

[26] An idealization, to be sure.

[27] According to the "US Cattle Supplies and Disposition" website, this is approximately the number of cattle in the US.

thinking about the situation seems to confer salience on the (by hypothesis, slight) possibility that he's wrong. After all, is he really sure that no cattle feed intended for Canadian consumption has mistakenly been shipped to US ranchers, or that runoff from feedlots over the border hasn't entered our water supply? This isn't a mere metaphysical possibility; things like this do happen, and there's no reason to think that this is a special case. And the stakes are certainly high, since the resulting disease is invariably fatal. In these circumstances, that is, the proposition "This burger is perfectly safe" has the status of a lottery proposition, since one can't help but think about—yet can't rule out—certain ("real," not "hyperbolic") counterpossibilities to the claim. And yet it seems normatively appropriate for the official to utter it as long as, after doing considerable research, he has excellent probabilistic grounds for his claim.

If we subsequently learn that he was covering up counterevidence about the probability of infection, or even that his true, highly probable, assertion was based on no research whatsoever, we'd consider him to have violated the norms of assertion, but in the case I described, we would not. Indeed, we may take his assertion to be normatively *required*, since the public may take any weaker claim (or silence) as suggesting (if not "implicating")[28] that the probability of there being infected cattle in the US beef supply is much greater than in fact it is.

For another example, I may say (truly) to my teenage son, who's been studying quite diligently for an exam, and in whom I want to show (and instill) confidence: "You'll do well on the test tomorrow." Presumably, I have excellent justification for my claim: he *has* been studying, and he's kept up with the readings and exercises all term. If he hadn't been doing this, I'd have no right to make my assertion; he has, however, and thus it seems that I do. But do I *know* he'll do well? Presumably not: all sorts of (non-hyperbolic) things could happen, as I am well aware; the teacher could be out of sorts, arbitrary; someone else in the class could do brilliantly, raising the curve. Am I thereby to be criticized for

[28] I do take seriously the worry (raised, for example, by DeRose (2002)) that it often seems unprincipled to invoke "implicature" in characterizing what's suggested rather than actually expressed.

making this assertion? I don't think so, as long as I have good evidence for it.[29]

There are other, more mundane examples, in which it seems that asserting is primarily doing the job of conveying information. Suppose I've just checked my e-mail, and announce to my colleagues, "Ed won't be coming to the meeting because he's sick." I can assert this, it seems, even if, at the time of the assertion, I'm aware of the fact that people sometimes make excuses. Or consider my reporting, late at night, "Henry still has his cough," upon hearing coughs coming from behind his closed bedroom door. This violates no norms, I submit, even if I realize that for all I know some of his friends may have come over after I went to bed, and it's one of them who is coughing. (I have no evidence that anyone has come over without my being aware of it, but sometimes people have.) One might object that what I'm *really* asserting in these situations is that Ed *says* he's sick, or that *it seems* as though Henry still has his cough. But in doing this I'd be suggesting (implicating?) that I have more doubt about the situations in question than I do.

Even when we stick to the classic cases used as evidence for the unassertability of propositions we don't know, the verdicts can be mixed. Indeed, it does seem wrong to assert, when waiting for the results of the lottery (or, even worse, when buying the ticket) "This ticket is a loser." But suppose, while all the time thinking about the lottery ticket in my pocket, I'm asked to fill out a form (or swear in court) about my estimated income for next year: under $50,000, $50–100,000, $100–200,000, or "don't know." And suppose my salary from my job is about $48,000/year. Can I, must I, answer "don't know?" If I answer "under $50,000," while being fully conscious of my lottery ticket, have I violated some norm? Indeed, can't I even explicitly follow up that assertion with "Of course, I just bought a ticket for the $5M Lotto (ha, ha)," without landing in the normative doghouse?

Consider the "car theft" cases. Suppose my spouse comes home, notices that my car isn't parked, as usual, somewhere on our block, and says: "I just heard a report on the radio of the annual car theft statistics for our neighborhood—and I see that you're home, but your car isn't anywhere on the block. Do you think we should call the police?" Can't I

[29] Sometimes the conditions can be even weaker. Consider "There are WMD's in Iraq?" This assertion, of course, turned out to be false, and there has been general outrage that US government policy was based on this premise. However, the outrage comes not because the administration didn't *know* this, but because they asserted it when there was so little evidence *for* it! If, however, there had been substantial evidence to support the claim, then it would have seemed appropriate to assert it, and criticism of US policy would have to be confined to its practical *conclusions* ("Therefore, let's go to war"), rather than the information-conveying premise on which they were based.

say: "No, don't worry, it hasn't been stolen; I just parked it a couple blocks away on Elm Street." (Wouldn't it be wrong *not* to say this—or to say "Come to think of it, I don't know where my car is!")[30]

Indeed, sometimes we find it normatively appropriate for people to make assertions when they explicitly *refuse* to represent themselves as knowing. For example, a recent LA Times article about an 11-year old female boxer used as its title her assertion "I'm Gonna Win. But You Never Know."[31] The boxer had never been beaten, and did go on to win the match, so her "I'm gonna win" was justified and true. But it neither was, nor was represented as, knowledge.

These scenarios were supposed to be cases in which S's asserting that p is normatively permissible (or required), even though, given the circumstances of the subject or the ascriber (or both), S can merely be said to have well-justified true belief, or well-justified false belief—but not knowledge—that p. Still, some might object that these cases, even if compelling, aren't really counterexamples to the thesis that knowledge is the norm of assertion. Sure, one might concede, people say things for a variety of reasons, or are interpreted as meaning a variety of things when they assert that P.[32] But this just means that sometimes

---

[30] Indeed, sometimes it's normatively incorrect to refuse to assert a proposition, or to claim that one doesn't have knowledge of that proposition, even if this in fact is true. Here's another example from (relatively) current events:

> "Over the past several months, Mr. Cheney has aggressively sought to tie foreign terrorists, specifically Al Qaeda, to Iraq. In September, a few days after Mr. Cheney said the government did not know whether Mr. Hussein had some connection to the Sept. 11, 2001 terrorist attacks, Mr. Bush all but contradicted him. Asked by reporters about Mr. Cheney's statement, the president replied, "No, we've had no evidence that Saddam Hussein was involved with September 11th." (Eric Schmitt, NY-Times, November 11, 2003)

Bush's statement, of course, does not (even "all but") contradict Cheney's: having no evidence that p does not imply that one does not know that not-p. A better characterization of what Bush was doing, it seems, is suggesting that it's inappropriate to *assert* that we don't know that Hussein wasn't involved in these attacks, even though it's true. (Perhaps this is because the statement has a pragmatic implicature, namely, that government officials have evidence of SH's involvement that may surpass ours. Indeed, perhaps it is this implicature which Bush's statement can be taken to contradict.)

[31] Kurt Streeter, La Times, July 13, 2005. Later in this section I'll explicitly discuss Moore's Paradox.

[32] For example, consider the comment made by Timothy Noah, in Slate on-line magazine (July 14, 2005): "As I explained yesterday, when a president or his spokesman expresses full confidence in a controversial aide, that's usually a sign that the aide is about to get canned. Yesterday, McClellan wouldn't comment at all about Rove, which suggested to me that Rove's resignation was not imminent.... When McClellan says, "The president is behind Karl Rove 100 percent," start looking for the ax to fall."

assertions (even those intended strictly to convey information, and taken literally) can be overridden by various social or pragmatic considerations.[33] Thus none of my examples, even if intuitively compelling, undercuts the claim that knowledge (and knowledge alone) is the norm of assertion.

In the BSE case, perhaps, the primary point of such utterances is to calm a nation's fears, and thus the intuition that it's normatively appropriate, or even required, to assert that the burger is safe is an intuition about when it is that pragmatic considerations override our epistemic norms of assertion. Similar considerations may apply to the car theft case I sketched above: here it wouldn't be right to alarm my spouse, who is imagining that if my car isn't parked on our block, then it's sure to have been stolen. One might argue, that is, that for pragmatic reasons, sometimes one can appropriately assert p as long as the chances of p's being false are slight, and (perhaps in addition) there would be bad consequences, or false implicatures, if one failed to assert p. Other pragmatic considerations may figure into our intuitions about other cases as well. In court, for example, adherence to a cautious, purely epistemic, standard for making assertions may be taken to signal insubordination, or lack of regret for one's actions.[34] And no doubt there are pragmatic reasons why the young boxer can't say that she knows she's going to win; this, presumably, would display hubris and alienate her fans. All these apparent counterexamples, however, can be regarded as cases in which non-epistemic, purely pragmatic, norms override our bedrock intuition that knowledge is the norm of assertion. It's only in these aberrant cases, that is, that it seems that the norm for assertion can be something as weak as justified true—or even justified false—belief.

The problem with this objection is that it seems that there are so *many* aberrant cases. This provides motivation to explore the possibility that the norms of assertion are *always* pragmatically determined: depending on one's circumstances and interests, one sometimes can be normatively correct in asserting that p only if one has justified belief

---

[33] See Williamson (2000, p. 256).

[34] For example, an article in the Washington Post (July 1, 2005) reports: ''A Michigan man who threatened to blow up his van near the White House two days before President Bush's second inauguration stunned a court yesterday by saying he couldn't promise that he wouldn't do it again....'' There's always a chance of anything, Your Honor,'' he said... The judge's jaw dropped. He pressed Timmers [the defendant] to be clear. ''The odds of that happening are 800 million billion to one,'' Timmers said. ''But I can't ever rule out anything completely, Sir.'' (Carol D. Leonnig) Thanks to Luka Yovetich for bringing this to my attention.

that p, other times, only if one's justified belief that p is also true, yet other times, as long as one has (mere) true belief that p—and, in certain cases, only if one knows that p. On this view, the pragmatic element in the evaluation of assertions attaches not to one's epistemic credentials (whether one has knowledge or justified belief), and not to one's state of mind (whether one has a bona-fide belief, or a mental state with a somewhat different functional role) but to the norms of assertion themselves. Sometimes, to assert p legitimately, we require a person to know that p—but more often, something lesser will suffice.

Indeed, even in cases in which we lack knowledge of propositions that seem *in*appropriate to assert, it's not clear that our lack of knowledge is what explains the impropriety. For example, consider the assertions of "lottery propositions" that appear to violate the norm of assertion: "My ticket is a loser," uttered when waiting for the lucky number to be drawn, or "My car hasn't been stolen," uttered after someone has reminded me that there are occasional thefts in this neighborhood. First, think about the circumstances in which it seems inappropriate to utter "My ticket is a loser" (when I've bought a ticket in a lottery in which the chances of winning are extremely low). Here, I'm often in a situation that's the result of some sort of irrationality: I've bought a *lottery ticket*, for God's sake, in a million-ticket lottery! What was I thinking?

There are a number of possibilities (that readily come to mind). First, it may be that I deluded myself into thinking that my chances were better because there were many fewer tickets for sale. Or perhaps I bought a ticket numbered with my birth date, figuring that this was my lucky number, and thus I have a pretty good chance of winning. In either case, I think (unjustifiably, to be sure) that my chances of winning are better than they are, and thus don't even believe that my ticket is *likely* to be a loser; this, it seems—and not that I don't have knowledge—explains why it wouldn't be appropriate to assert that it is.

Alternatively, I may have bought the ticket for the entertainment value of being part of a crowd that follows the build-up of funds in the pot, and gets excited when the winner is about to be announced. This is often the motive of people who play the lottery. But the entertainment value of my ticket depends on my thinking that it might be a winner, and so it would be pragmatically inappropriate for me to assert, under those conditions, that it's not. Finally, it may be that I bought my ticket on a whim (or even as a matter of habit) without thinking too much about my chances of winning. But even now, when I'm faced with a sobering recognition of my million-to-one odds, it seems inappropriate to assert that my ticket is a loser. This, however, could be explained as an attempt by me to honor an earlier

commitment; to refuse to repudiate an action no matter how ill-advised it seems in the light of day. Lotteries, in short, bring out the irrational in their participants (in one way or another), and so one can expect that the norms for asserting lottery propositions have pragmatic features that reflect (or attempt to hide) this fact.

Similarly, the possibility that my car parked on Elm Street has been stolen—or, the possibility that I (as a healthy individual with no family history of heart disease) will have a fatal heart attack within the year—will prevent me, on most high standard views, from having knowledge. But here, too, I want to deny that my lack of knowledge explains why assertions to the contrary seem inappropriate. In these cases, like the lottery cases just discussed, it seems that the norms prohibiting such assertions are pragmatic. But, while in the lottery cases, my motivation is to represent my chances as better than they are, in these cases, my motivation is to represent my chances as *worse* than they are—and thus avert the evil eye, or attention from any other malign forces that may frown on good fortune. That is, there are often *superstitious* reasons for people to find it inappropriate to make certain assertions, no matter how well justified they are.

To be sure, these cases may give support to *low* standard invariantist theories, since if the inappropriateness of making these assertions is due to factors other than the speakers' lack of knowledge, a low standard theorist can explain it in the same way. But, as I've stressed, though I'll give considerations in favor of a high standard theory in the final section, my main motivation is to dispute the claim that knowledge is the norm of assertion—no matter what the consequences. Indeed, I'll argue next for another claim that can aid a low standard invariantist theory, namely, that there are some cases in which it seems that knowledge may not even be *sufficient* for assertion.[35]

Consider, for example, the variations on DeRose's well-known "Hannah and the bank hours" case. When the stakes are extremely high (e.g. when Hannah needs the money to pay ransom to her child's kidnappers, as in Russell (2004)), she can't rely on her knowledge that the bank was open on Saturday two weeks ago (banks change their hours), or even on the sign posted on the bank's door (the sign may be

[35] This would explain, for example, why Hannah, John, and Mary must keep checking, even though, on a low standards view, they have knowledge of the propositions ("The bank is open on "Saturday," "The plane stops in Chicago") in question.

out of date), but has to go inside and talk to a teller at the window. But will even this be enough? One might think not: Hannah might want to talk to the manager. Still, even if she has done this (and has also somehow managed to determine that the bank manager is not a disgruntled employee, bent on passing on misinformation to the customers),[36] she might want to call again when she gets home (if she's like at least some of us), just to make sure.

However, this need not be because her belief about the bank's hours isn't knowledge, no matter how high the standards for knowledge are (short of Cartesian skepticism). She'd be judged as normatively deficient, I suspect, if she didn't keep checking, obsessively, to make sure that the hours haven't changed: after all, it *is* her child! That is, sometimes obsessive, neurotic reinforcement of the knowledge one has is the only acceptable response to a situation, since it's the only way for a person to take action in that situation—or show others that they care.[37]

There are other cases in which it seems normatively inappropriate to assert, or act on, a proposition, even if one counts as knowing it—again, on almost *any* view about what the standards for knowledge are. If I can't find the expensive ring you gave me, and have overwhelming reason to believe that it's been stolen (I've heard that there were jewel thieves in the neighborhood, and saw a masked figure slip out of my bedroom window as I was walking up the steps to my house), I may refuse to assert this, and continue to look for it—not

[36] I'm grateful to an anonymous PPR referee for pointing out that, on the high standards invariantist view, Hannah has to do a lot more than check with the manager—e.g., make sure that the manager is not a disgruntled employee—to rule out counterpossibilities even to the thesis that the bank's policy is to be open on Saturday. Further, if Hannah can rule out such counterpossibilities, *all* she can know is a proposition about bank policy; she can't know that it *will* be open on Saturday unless she can rule out the possibility that the bank will burn down on Friday night. But even if all Hannah can know is a proposition about the bank's policy, my contention is that it still may be normatively appropriate for her to keep (obsessively) double-checking about *that*.

Perhaps what Hannah (along with the other subjects I'll soon be discussing) is doing, in these cases, is thinking of certain mere "hyperbolic" counterpossibilities which cannot be ruled out—though ruling them out is in fact unnecessary for knowledge, even on high standard views.

[37] As the caption of a recent Bruce Eric Kaplan cartoon puts it, "But if you don't cry over spilt milk, everyone thinks you're a cold bitch." (LA Weekly, July 29-August 4, 2005).

because I don't *know* that it's been stolen, but because I want to show you how important it was to me.[38] My suggestion, in short, is that there's a gap between knowing and acknowledging—on almost any account of the standards for knowledge—which is best explained by a pragmatic theory (not of knowledge, but) of assertion itself.

Even if I'm correct, however, any high standard view would seem to have a further problem, namely, the threat of (a version of) Moore's Paradox. Just as there seems to be something odd or paradoxical about my saying 'P, but I don't believe that P', even though that statement may be true (Moore's example), there seems to be something almost as odd or paradoxical in saying 'P, but I don't know that P'. A promising way of explaining this oddity is by taking knowledge to be the norm of assertion; my asserting P, on this view, is only normatively appropriate if I know (or represent myself as knowing) that P, so the infelicity of the Moorean statement can be explained as a matter of my implicating, and simultaneously denying, knowledge that P.[39] No such explanation, however, is available on my view.

It's not clear to me, however, that the "knowledge" version of Moore's Paradox *is* as paradoxical as it may at first seem. Some instances of it, to be sure, seem quite problematic. For example, in discussing this question, Jason Stanley (2004 ms. note 17) quotes Moore (*Commonplace Book*) as remarking, "Dogs bark, but I don't know that they do" [is absurd]…because by asserting p positively you imply, though you don't assert, that you know that p." Stanley takes this as evidence that knowledge is the (constitutive) norm of assertion. But

---

[38] Here, too, one might wonder whether this proposition will count as knowledge on the high standards view. But if it falls short, I hope I've made it plausible that this is not the problem with making the assertion.

There are other cases, it seems, for which this is true as well. Suppose I have overwhelming reason to believe that I—or you—have failed at some important project, or that I—or you—have some terminal disease. There are many reasons, it seems, why it may be inappropriate to assert these things, no matter what one's standards are for knowledge. To be sure, those who take knowledge to be the norm of assertion can argue that it can be overridden by *other* norms, and thus that these cases can be accommodated by their view. But, once again, there seem to be *many* such cases, and a pragmatic theory of assertion (rather than knowledge) would encourage the investigation of the interesting similarities and differences among them.

[39] See Stanley (2004), ms. p. 23 and DeRose (2002), p. 181. Also see Williamson (2000, p. 253), who also suggests that the oddity of the Moore statement can be explained by taking knowledge to be the norm of assertion. For then "to have warrant to assert… 'A and I do not know that A' is to know that A and one does not know that A. But one cannot know [this]. One knows [it] only if one knows each conjunct, and therefore knows that A (the first conjunct)… yet one knows the conjunction only if it is true… so only if one does not know that A (the second conjunct); thus the assumption that one knows [the conjunction] yields a contradiction." Jonathan Adler provides a similar diagnosis in his (2002), and goes on to develop a related treatment of the lottery paradox (Chapter 7).

this, I think, moves too fast. This statement is certainly odd, but another explanation is that, in this particular case, we think that anyone who takes it upon themselves to assert that dogs bark must know that they do, because they've acquired the belief via a good, knowledge-conferring, method for knowing, i.e., perception (and, perhaps, memory). They believe that dogs bark, that is, because they've heard them. But in cases in which this condition is not met, the appearance of oddity diminishes. For example, does the assertion "Hyenas laugh, but (come to think of it) I don't know that they do" seem equally odd? What one is doing here may be merely a matter of conveying conventional wisdom, while recognizing at the same time that this is all one can do.[40] But we assert things all the time, it seems, under these conditions, without violating any norms.

Still, even if the second assertion is less odd than the first, there remains an air of oddity about it, and there ought to be a reasonable explanation for this. My suggestion is that the Moorean assertions involving claims not based on knowledge-conferring methods like perception appear paradoxical primarily because of their association with their uncontroversial siblings, Moorean assertions about *belief*. Nearly everyone agrees that it's extremely odd, and perhaps even incoherent, to say "P, but I don't believe that P"—and there are many explanations of this oddity.[41] As long as asserting P commits one to believing (rather than knowing) P, then I can avail myself of these explanations. However, if one gets used to thinking of knowledge as a relatively rare achievement—a way of thinking made easier by the recognition that it's often normatively appropriate to *assert* that p without knowing p—then these Moore-statements, in certain circumstances, may seem acceptable, and just as modest (and perhaps endearing) as the statement of the young boxer who asserted "I'm gonna win, but you never know!"

---

[40] One might wonder whether this claim amounts merely to my asserting that hyenas laugh, and then taking it back when I realize that I don't in fact know this—thereby *affirming* the view that knowledge is the norm of assertion. I'm not sure what evidence would count definitively for my position (that is, that it's often normatively appropriate to make assertions that convey (or represent oneself as believing) conventional wisdom), but it seems that normally people aren't faulted for asserting things on these grounds. It also seems odd, as it shouldn't on the alternative view, to say "Hyenas laugh—but I guess I shouldn't have said that" upon realizing that one doesn't *know* that they laugh. Should my utterance be construed as shorthand for "It's conventional wisdom that hyenas laugh?" No, for the reasons given in the discussion of other examples in the text, namely, that this would convey some (practical) doubt about the conventional wisdom that we're not taking the speaker to have. Thanks to Jonathan Adler for pressing me to say more about this case.

[41] See, for example, Shoemaker (1996), Adler (2002), and many others.

IV. Knowledge and Practical Reasoning

But isn't it normatively appropriate to rely only on what we know as premises in our practical reasoning? Sometimes, I suppose, but very often not. If I don't have much money, it seems quite reasonable to look into vacations on the Jersey shore (rather than Africa) next year, despite my preference for a trip to Africa—even if I'm thinking about that lottery ticket in my pocket that would net me $500,000 if it were a winner. And it also seems quite all right for me to offer you a ride home if you walk with me to Elm Street, even if I remark that of course my car might have been stolen, because these days you never know. That is, even if it's quite salient to me that the premises of my reasoning entail lottery propositions that I don't know, it often seems normatively OK to use them in my reasoning (and unrealistic or curmudgeonly to hold back). In short, the norms of practical reasoning, like the norms of assertion, seem to be a diverse and complicated lot, involving a variety of practical considerations.

Similar considerations apply not only when the counterpossibilities to one's premises are salient, but when, in addition, the stakes in the deliberation are high. Suppose I'm deciding whether to take a job in Chicago or one in Madison, Wisconsin. I've made a meticulous cost-benefit analysis of my options, and have come to the conclusion that, though each has different pluses and minuses, their expected utilities are roughly the same. Except, that is, that I think the possibility of a terrorist attack on Chicago (or, more precisely, an attack near my would-be office) in the next 10 years—though slight—is greater than the possibility of a terrorist attack on Madison. I realize that I don't know I won't get killed in a terrorist attack if I move to Chicago (whereas the possibility that I'll get killed in a terrorist attack if I move to Madison is much more "metaphysical", and, like the Cartesian skeptical hypothesis, would not be taken to undermine my knowledge). Is it rational for me to take this into consideration when I make my decision? Many people would say "no."[42]

There may be other reasons, however, to take knowledge to be the norm of practical reasoning. Hawthorne, Stanley and Williamson discuss a number of cases that exhibit faulty practical reasoning, and argue that the faultiness can't be explained unless we take this to be so. Suppose, for example, that you've been offered 1c for your $1 ticket in a 10,000 ticket lottery with a $5000 prize, and reason as follows:

[42] I'm grateful to Ned Block for discussion of these points.

(1) I'll lose the lottery.

(2) If I keep the ticket, I'll get nothing; if I sell it I'll get 1c.

Therefore, I ought to sell the ticket.[43]

This does indeed seem like bad reasoning, but why? Hawthorne argues that the best explanation of why this arguments seems so bad is that we take it to be normatively appropriate to use only premises that one *knows* in one's practical reasoning. On this assumption, the fallacy in these arguments would be explained by noting that the first premise is something we can't know. This explanation is available both to contextualists and to "sensitive invariantists" like Hawthorne. Of course, it is also available to "high standard" invariantist views, like the one I've been discussing, as well.

But there is a more satisfying explanation, I contend, of what's wrong with these arguments that does not require the premises to be known, rather than justifiably believed. The reason, I suggest, that this piece of practical reasoning is intuitively bad is not because I don't know the first premise—the claim that I'll lose the lottery—but because I don't have justified belief in the *second* premise. In at least some circumstances, it's just not true (and I wouldn't be justified in thinking) that if I keep the ticket I'll get nothing. In many cases (as discussed in the previous section) I've bought the ticket, at least partially, for that *frisson* of excitement that's available only to ticket holders when the winner is announced. And in cases in which the anticipated excitement isn't a factor, I *still* may lose something in selling the ticket for 1c, namely, the conviction that I'm committed to the choices I've previously made (again, see the above discussion). If these benefits are worth more than 1c, then in selling I'd sustain a loss.

Even if we think only of the monetary pros and cons of selling my ticket, it's not clear that Hawthorne has the right diagnosis. Suppose that the prize in the lottery is not $5000, but $50,000, and the number of tickets, as in Hawthorne's example, is 10,000. And suppose I'm offered $5 for my ticket. Here the prize is greater (and the expected utility of selling vs. keeping is the same)—but it seems (at least to me) less crazy to sell. After all, I'd be getting $5, which is substantially more "better than nothing" than 1c!

Hawthorne has us consider another case of practical reasoning which he takes to be seems even more "intuitively awful"[44] than the

[43] See Hawthorne (2004), p.174.

[44] Hawthorne's term for *all* these arguments (2004, p. 174).

first. In this situation, you're offered a 1c ticket for a lottery like the first, and reason:

(1) I won't have enough money to go to Africa next year.

(2) So, if I buy a lottery ticket I'll lose.

Therefore, I shouldn't buy a lottery ticket.

This argument does seem intuitively awful. But the reason, once again, does not seem to be that the first premise isn't known. A better explanation may be that the way the second premise is worded suggests that there's some sort of strange causal relation involved; that my not having enough money to go to Africa next year somehow makes the ticket I'm being offered a loser. However, even if this suggestion is avoided by rewording the second premise to read, say, "The ticket I'm being offered is a loser," the argument still seems worse than the first. But here, I suggest, the problem is that practical reasoning standardly takes the form of means-ends calculation, and while the first argument conforms (more or less) to this standard, the second one does not: thus its startling oddity. Bringing the second argument into conformity with the standard would render it as (something like):

(1) I want to (do what it takes to) have enough money to go to Africa next year.

(2) I'll have enough money to go to Africa next year only if I buy a winning lottery ticket.

(3) The ticket I'm being offered is a loser

Therefore, I should not buy the ticket.

This piece of reasoning, however, seems eminently *rational*; isn't it almost always irrational to buy lottery tickets, given the odds (unless, that is, there are other reasons to buy lottery tickets besides the chance they'll give you of making money)? To be sure, this argument doesn't reflect the deduction of "This ticket is a loser" from "I won't have enough money to go to Africa next year" that Hawthorne wants to highlight as odd. But, once again, people rarely reason like this in determining what to do.

Hawthorne has a third example of intuitively awful reasoning that shares its structure with the second: Suppose you are offered life insurance, he proposes, and reason:

(1) I'll be going to Blackpool next year.

(2) So, I won't die beforehand.

Therefore, I ought to wait until next year to buy insurance.[45]

Now it does seem perverse to think of using "I'm going to Blackpool next year" as reason to refuse insurance, in part because, like the second argument, it violates the classic form of means-ends reasoning. Here, too, the relevant premises seem to be "I won't die before next year" and "It's only worth buying life insurance now if there's a chance I'll die before next year," yielding the conclusion that I shouldn't now buy life insurance. But here, too, once we put things this way, the argument seems less intuitively awful. Suppose, that is, I substitute for the premise "I won't die beforehand" the claim "I'm young, healthy, and will stay this way for the next year." Is it irrational to refrain from buying life insurance on the grounds of *this* premise? If one is really justified in believing it, I'd say not, even if one doesn't know it. After all, it's not always irrational for (say) healthy 25-year olds to put off buying insurance until a later time—even if they can't *know* they won't die within the year. At some point the probability that one will die becomes sufficiently great that it does seems irrational to put off buying insurance—but not always. In short, it seems that there are better explanations for the intuitive awfulness of these pieces of practical reasoning than that some lottery proposition, or a proposition that entails it, is used as a premise in those inferences without being known.

But here, too, this alternative explanation of the impropriety of these arguments could be invoked by a *low* standard theorist. Though it raises a challenge to the claim that knowledge is the norm of practical reasoning (my primary goal), it doesn't particularly give support to the high standard view. So again one might ask why I've expressed even a tentative preference for a high standard invariantist theory.[46] Why not endorse a *low* standard theory, or remain neutral between the two?

V. High vs. Low Standard Theories

On a low standard view (at least arguably),[47] Hannah could know that the bank will be open on Saturday as long as she reads a sign

[45] Hawthorne (2004), p. 175.

[46] I'm grateful to an anonymous PPR referee for prompting me to clarify this point.

[47] Here, too, any final verdict will depend on exactly what the standards in question are.

posted on the door that says so—no matter how important it is that her check gets deposited before Monday—and Smith, John, and Mary could know that their plane stops in Chicago if that's what it says on the recently posted schedule—no matter what the relative importance is of their stopping in Chicago, or how salient the counterpossibilities are for them, or for anyone trying to determine whether or not they know. And as long as I don't have (or could easily get) specific evidence to the contrary, I could know not just that my car is parked on Elm Street, but that it hasn't been stolen, and that I won't have a fatal heart attack within the week.

On such a view, of course, I would also know, even before the drawing, that my lottery ticket is a loser, and this seems intuitively unacceptable to most "pragmatic" theorists—and to others as well. Further, many argue, if it's unacceptable to think we know that our ticket is a loser, then it's unacceptable to think we know that our car hasn't been stolen or that we won't have a heart attack within the week—since these "lottery propositions" are all in the same boat.[48]

But, just as with high standard theories, one can argue that these intuitions get particular force from the assumption that knowledge is the norm of assertion and practical reason. A successful challenge to this assumption, therefore, would also enhance the prospects of a low standard theory, since one need not be concerned that sometimes it seems inappropriate to assert, or act upon, propositions that we can be said to know. The decision, thus, may ultimately depend upon which sort of theory gives the most intuitively plausible verdict about the *scope* of our knowledge of the world.

There are at least some intuitive considerations in favor of a high standard theory: it would honor the intuition, voiced by many (and which I myself find compelling), that once we're made aware of the counterpossibilities we'd been ignoring in cases like the "bank" and "airport" case, we agree that we don't (and didn't) in fact have knowledge—and that this verdict tends to stick. Likewise for claims about lotteries, fatal heart attacks, and car thefts. My own view, indeed, is that it's not intuitively implausible to hold that many *other* of our "ordinary" claims about the world, especially about the future and unobserved events, cannot be known (though they can, in many cases, be justifiably believed).

---

[48] Again, see Adler (2005), for dispute.

The greatest challenge for a high standard view, I contend, is to make a principled distinction between real and hyperbolic counterpossibilities that can prevent the slide into Cartesian skepticism, and can permit us knowledge of at least some "ordinary" propositions, based on perception, memory (and even, in some cases, on testimony and inference from well-established laws). If this can be done, however, then there is something intuitively appealing about a view which counsels modesty in our claims to knowledge. To be sure, there will be many ordinary propositions which I won't be able to claim I know; indeed, a view like this may be a species of what Vogel (1990a) calls "semi-skepticism."[49] But if, on reflection, I recognize that there are (real, not hyperbolic) counterpossibilties to those propositions that I can't rule out, it seems intuitively plausible to think that maybe I *don't* know them after all. And if I can assert them, and act upon them, in ordinary situations, then this recognition, though (I hope) sobering, will have little practical effect.

Still, we will no doubt continue to *claim* to have knowledge in these situations, and on a high standard view these claims will be strictly speaking *false*. Thus one might object to a view that makes our ordinary knowledge claims subject to such massive error. But, once again, if knowledge is not the norm for assertion or practical reason, then such massive "error" would have few practical consequences (except, perhaps, as Hume noted, when reflection on our condition prompts us to be in general less credulous). We could thus accept this result with equanimity, and go happily back to the backgammon table.

But a view like this, one might argue, robs knowledge of its *importance*. This is true. If knowledge is practically insignificant, then the most interesting work in epistemology will be the exploration of what it is to have justified (or reliably produced or responsible or virtuous) belief. Knowledge (in contrast to these other notions), in short, will play a much reduced role in epistemological inquiry as well as daily life. Indeed—though this is more the legacy of Gettier than of Hume—this had been the consensus in contemporary epistemology until recently; theorists tended to focus less on knowledge, and more on notions such as justification and epistemic responsibility.

---

[49] On this view, for example, (Vogel, 1990a, pp. 20–21), I can't know that Bush is president if I can't rule out the possibility that he has just had a fatal heart attack, and I can't know that there's a luncheonette near my office if I can't rule out the possibility that a fire has just burned it down. But (contrary to Vogel) it seems to me *correct* that I don't know these things, since on reflection I realize that there *are* real possibilities that I can't rule out. And, once again, if there are situations in which it's perfectly appropriate for me to assert or act on these propositions, the intuitions that support the claim that we *do* know these things have much less force.

One important motivation, it seems, for both contextualism and "sensitive" or "interest-relative" invariantism is the desire to reverse this trend, and give knowledge itself a bigger, more central, role in epistemology.[50] The considerations adduced in this paper, then, can be seen as raising the question: If knowledge is *not* the norm of assertion or practical reason, then why should we care?[51]

References

Adler, J. (2002). *Belief's Own Ethics*. Cambridge, MA: MIT Press.

—— (2005). "Reliabilist Justification (or Knowledge) as a Good Truth Ratio," *Pacific Philosophical Quarterly*, Vol. 86, No. 2.

Austin, J. L. (1946). "Other Minds." *Proceedings of the Aristotelian Society*, supp. 20.

Brueckner, A. (2002). "Williamson on the Primeness of Knowing." *Analysis 62.3*.

——. (2005). Review of *Knowledge and Lotteries*. *Mind* 114.

——. (forthcoming). "Evidentialism and Pragmatics Reconsidered."

Cohen, S. (1988). "How to be a Fallibilist." *Philosophical Perspectives* 2.

—— (1999). "Contextualism, Skepticism and the Structure of Reasons." *Philosophical Perspectives* 13.

——. (2004). "Knowledge, Assertion, and Practical Reasoning." in *Sosa and Villanueva*.

Conee, E. and Feldman, R. 2004. *Evidentialism*. New York: OUP.

DeRose, K. (1992). "Contextualism and Knowledge Attributions," in *Philosophy and Phenomenological Research* 52.

——. (2002). "Assertion, Knowledge and Context." *Philosophical Review* Vol. 111, No. 2.

Dretske, F. (2005). "The Case Against Closure," in *Steup and Sosa*.

Fantl, J. and McGrath, M. (2002). "Evidence, Pragmatics and Justification." *Philosophical Review*, Vol. 111, No. 1.

Harman, G. and Sherman, B. (2004). "Knowledge, Assumptions, Lotteries," in *Sosa and Villanueva*.

Hawthorne, J. (2004). *Knowledge and Lotteries*. New York: OUP.

[50] See again DeRose (2002, p. 182). Retaining the importance of knowledge is motivation as well for Williamson's argument that knowledge, rather than anything falling short of it, plays an essential role in *explanation*. See Brueckner (2002), Kaplan (2003), and Jackson (2002) for (in my view, quite convincing) criticisms of this argument.

[51] I'm grateful to Jonathan Adler and Anthony Brueckner for many helpful comments and suggestions. I'm also indebted to an anonymous referee for PPR whose comments led to substantive revisions in both the arguments and structure of the paper.

——. (2005). "The Case for Closure," in *Steup and Sosa*.
Jackson, F. (2002). "Critical Notice of *Knowledge and its Limits* by T. Williamson." *AJP* Vol. 80, No. 4.
Kaplan, M. (2003). "Who Cares What You Know?" *Philosophical Quarterly* 83.
Lewis, D. (1996). "Elusive Knowledge." in *Australasian Journal of Philosophy* 74.
McGrath, M. (2004). Review of John Hawthorne's *Knowledge and Lotteries*, for the Notre Dame On-Line Philosophical Reviews.
Russell, B. (2004). "How to be an Anti-Skeptic and a Noncontextualist." *Erkenntnis* 61.
Shoemaker, S. (1996). "*Moore's Paradox and Self-Knowledge", in Shoemaker, The First Person Perspective and Other Essays*. Cambridge: CUP.
Sosa, E. and Villanueva, E. (2004). *Philosophical Issues 14*. Boston: Blackwell Publishing.
Stanley, J. (2004). "Context, Interest-Relativity, and Knowledge" ms.
Unger, P. (1975). *Ignorance: The Case for Skepticism*. Oxford: OUP.
—— (1984). *Philosophical Relativity*. Minneapolis: University of Minnesota Press.
Vogel, J. (1990a). "Are There Counterexamples to the Closure Principle?" In M. Roth and D. Ross, eds, *Doubting (Philosophical Studies Series 48)*. Dordrecht: Kluwer.
——. (1990b). "Cartesian Skepticism and Inference to the Best Explanation". *Journal of Philosophy*, Vol. LXXXVII, No. 11.
——. (2004a). "Skeptical Arguments," in *Sosa and Villanueva*.
——. (2004b). "Speaking of Knowledge," in *Sosa and Villanueva*.
——. (2005). "The Refutation of Skepticism." In M. Steup and E. Sosa, eds, *Contemporary Debates in Epistemology*. Malden, MA: Blackwell Publishing.
Weatherson, B. (2005). "Can we do without pragmatic encroachment?" ms.
Williamson, T. (2000). *Knowledge and its Limits*. Oxford: OUP.